



department of computer science
faculty of sciences

Internet Programming

Client-Server Programming

Introduction

Introduction

- Now that you know how to make computers communicate, you can write **any** distributed application!
 - Decide which machine does what
 - Identify when computers need to communicate and what they should tell each other
 - Program it!

THE END

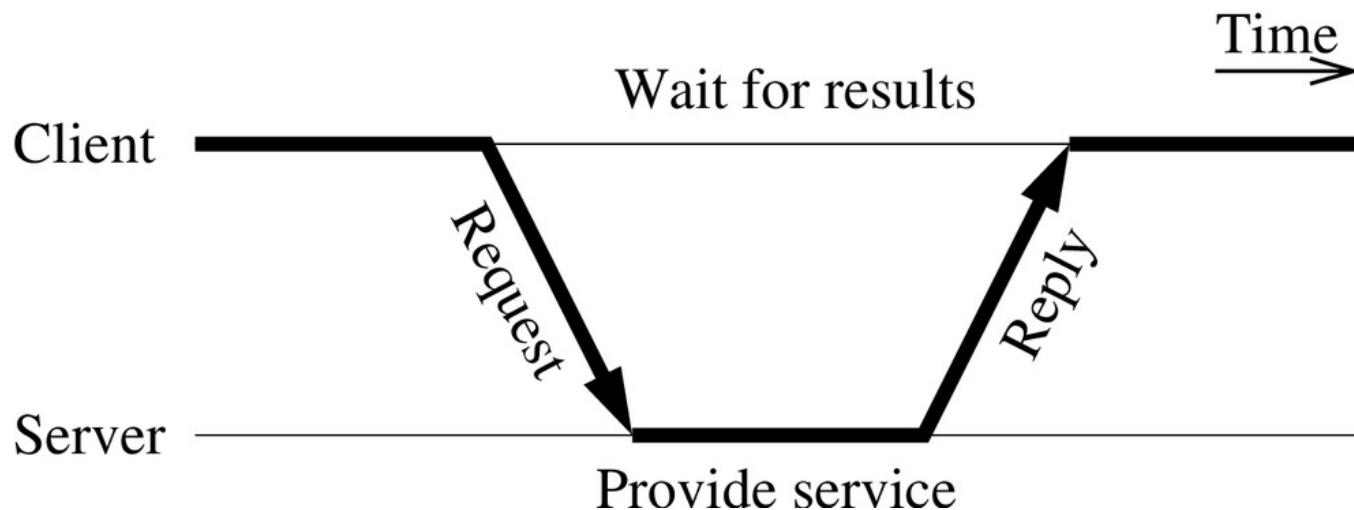
(well, **almost...**)

...except that

- You will have to deal with many issues:
 - How to **structure** your program
 - Define an **application protocol**
 - That is powerful enough for all your needs
 - Yet efficiently implemented
 - Deal with machines of **different architectures**
 - They may represent data differently
 - **Locate machines**
 - “Which machine implements task X?”
 - etc.
- You are in need of a middleware
- A **middleware** is a piece of software in charge of these issues
 - You program the application code
 - The middleware takes care of distribution issues (at least some of them)

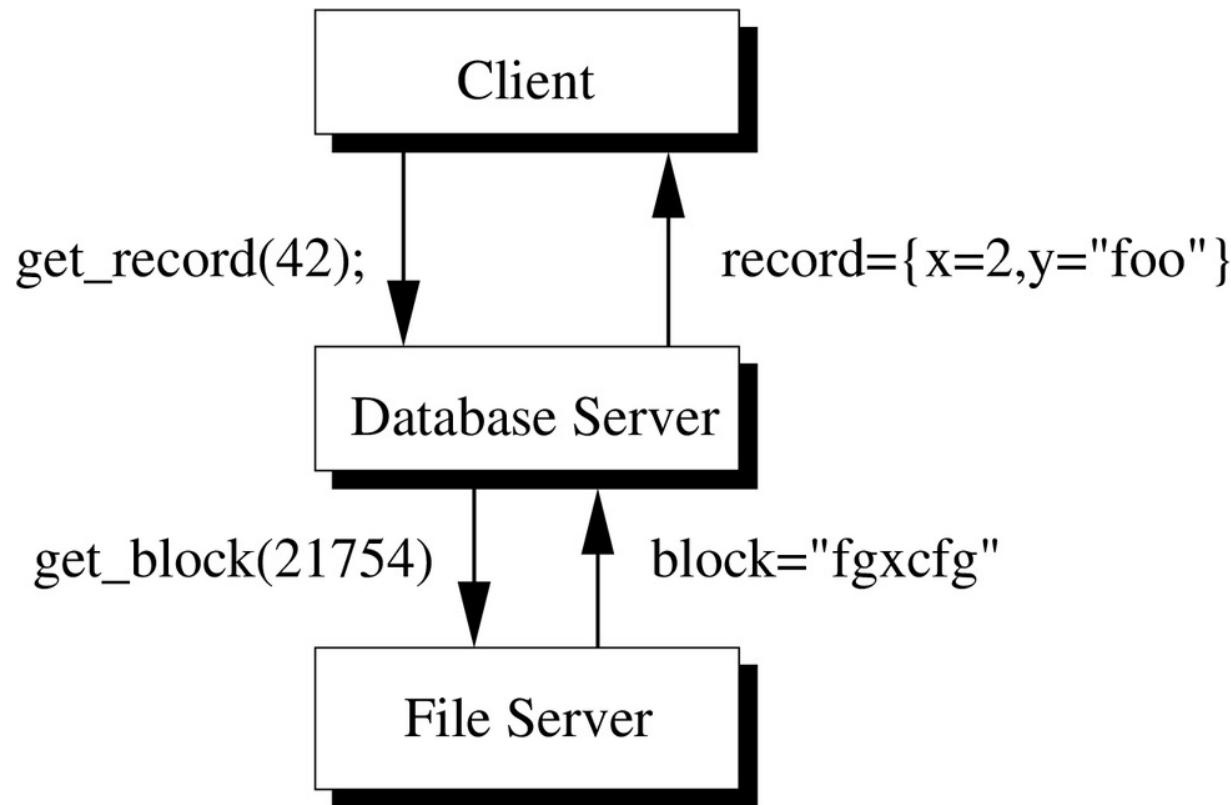
The Client-Server Model

- This is the most used model for organizing distributed applications
- **Servers** implement specific services
 - e.g., a file system, a database service
- **Clients** request services from the servers, and wait for the response before continuing



Chained Client-Server Interactions

- A server can itself be a client to another server:



- You just have to be careful about loops (A→B→C→A)

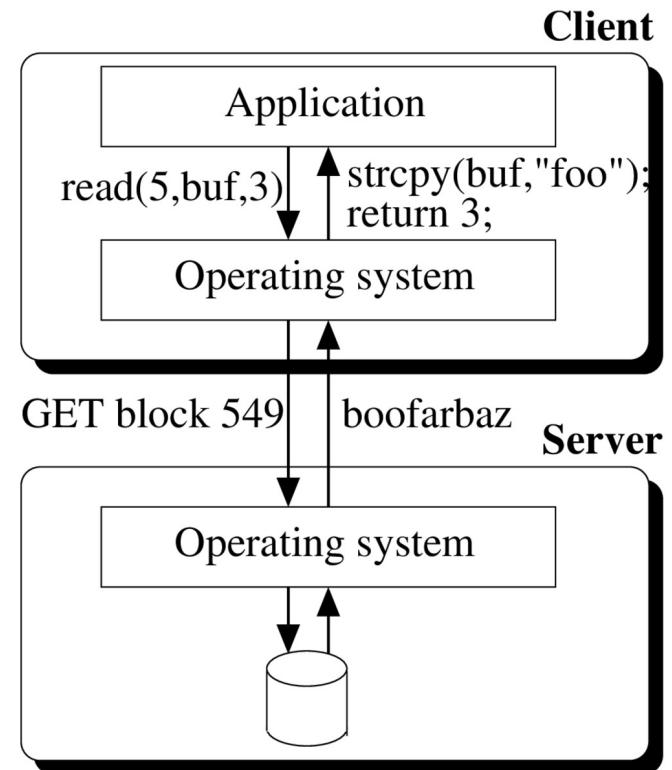
Why use the Client-Server Model?

- A client-server application (usually) **does not run faster** than a centralized application
 - The client is waiting while the server works
 - Additional delays due to communication
- But it has several advantages!
- **QUESTION:** What are they? ☺
 - Benefitting from **specialized resources**
 - Fast CPU, lots of memory, etc.
 - Special device attached (printer, disk array, GPS clock, etc.)
 - **Splitting up an application**
 - If it is too big to fit in one computer (memory space, disk, etc.)
 - **Sharing information** between multiple clients
 - A file server allows file sharing between multiple clients
 - Same for a database server (data sharing), etc.

Example

□ Distributed File Systems

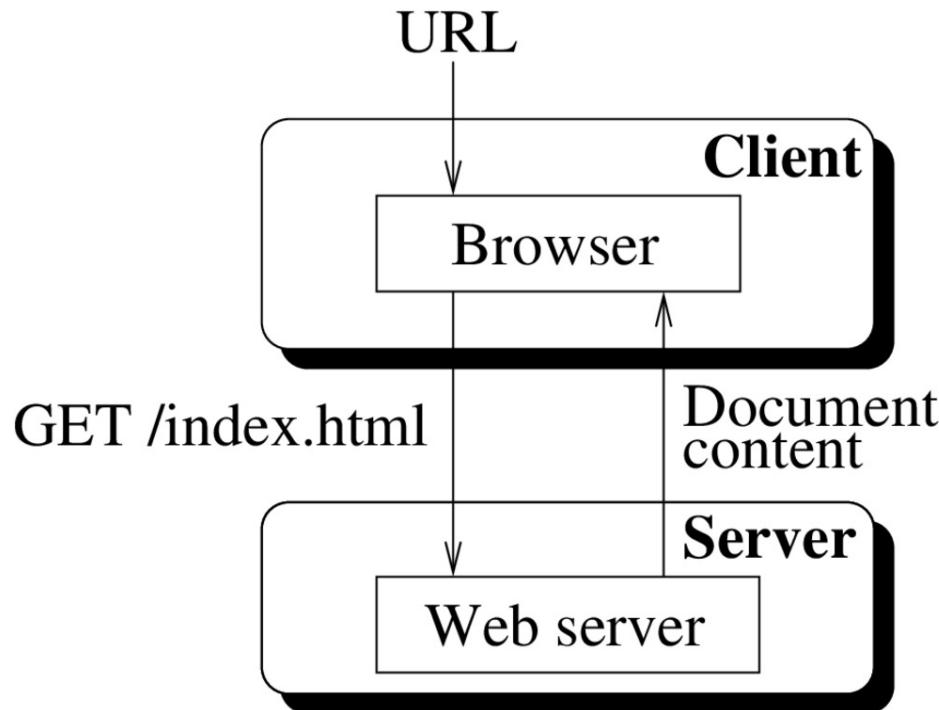
- The operating system contains specialized code
 - Convert file-related system calls to requests to the file server
 - Convert server replies into system call return values
-
- The whole system is built specifically for one given application



Another Example

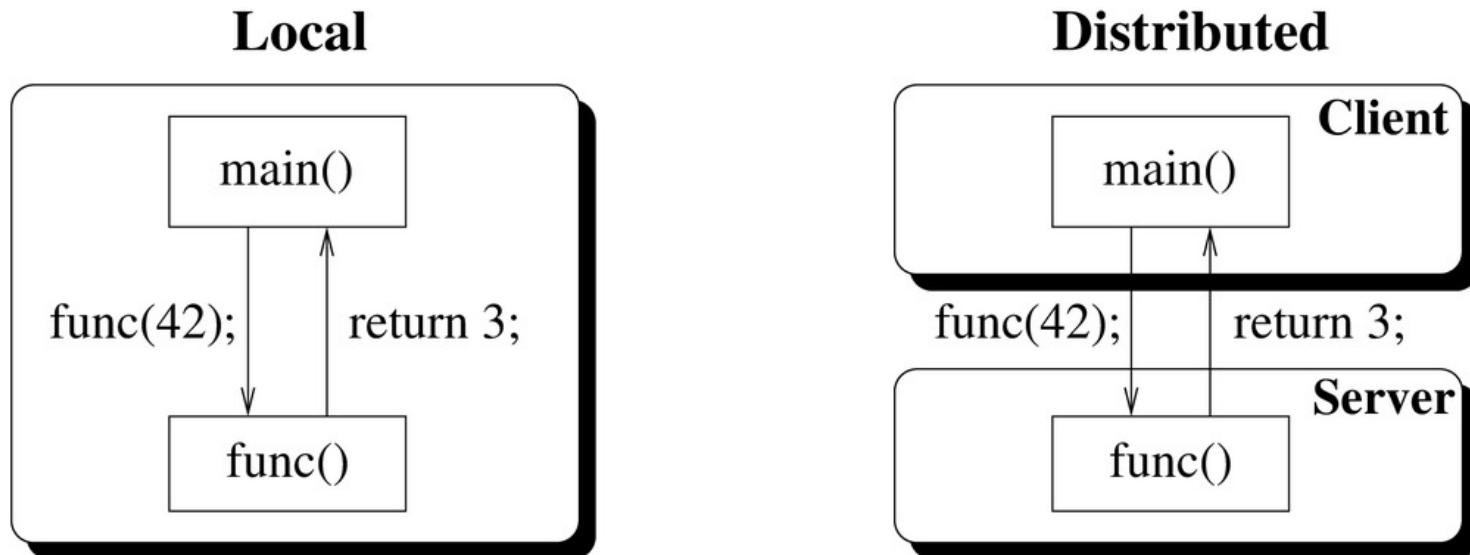
□ The World-Wide Web

- A specialized client-server protocol has been defined: HTTP



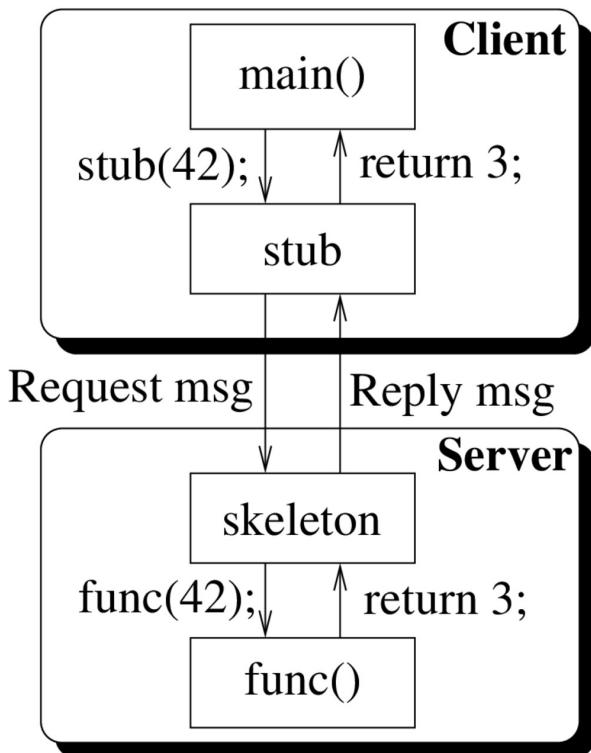
Remote Procedure Calls [1/2]

- There already exists a way to represent a task in a local application:
procedures (or functions)
- Let's extend the model to **remote procedures**



Remote Procedure Calls [2/2]

- You need to convert invocations into network messages and vice-versa
 - A **stub** is a function with the same interface as `func()`: It converts function calls into network requests, and network responses into function returns
 - A **skeleton** converts network requests into function calls and function responses into network replies
- An RPC system is used to generate the stub and the skeleton (more or less) automatically
 - Based on a description of the interface of `func()`

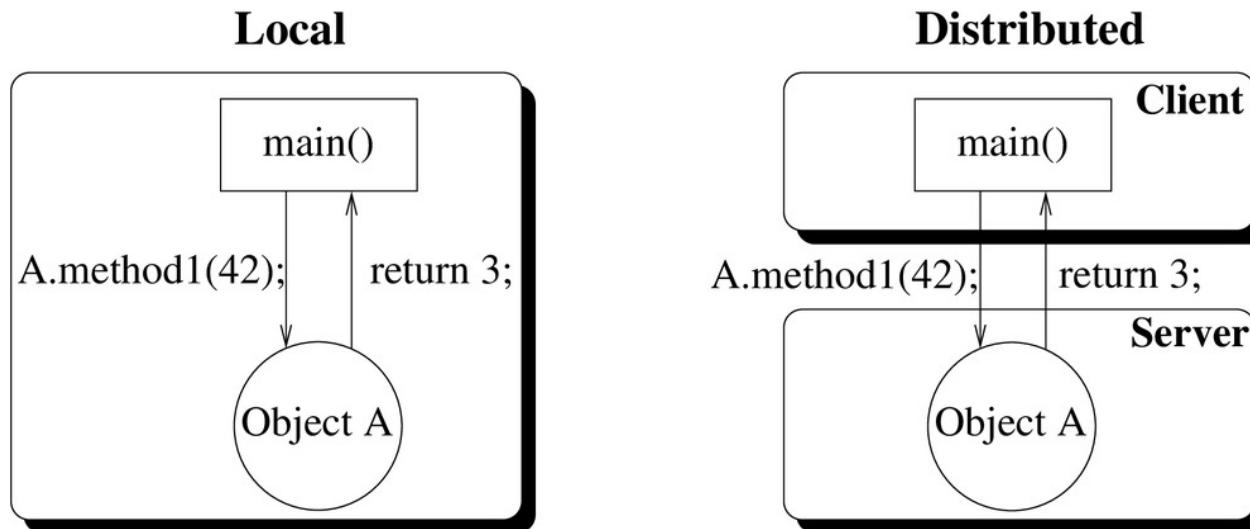


Limitations of RPC

- Clients and servers do not share the same **address space**
- Contrary to non-distributed programs, clients and servers:
 - Do not share **global variables**
 - Do not share **file descriptors**
 - Therefore the server cannot directly access a file opened by the client
 - Cannot use **pointers** as function parameters
 - Because the server will not be able to follow such pointers
- This sets constraints on which parts of a program you can separate and run as a server
 - The server must have a clear **interface** (a set of function prototypes)
 - The server can have **internal data**, but clients **cannot access** them
 - The client can have **internal data**, but the server **cannot access** them

Remote Method Invocation

- The equivalent to RPC in the object-oriented world is **RMI (Remote Method Invocation)**



- Like in RPC, you must have **stubs** and **skeletons**
- There are several Remote Method Invocation systems:
 - Sun RMI (entirely in Java), Corba (language independent), etc.

Using Sun RPC

Introduction

- Sun RPC is a Remote Procedure Call system
 - Officially called ONC-RPC
 - “Open Network Computing Remote Procedure Call”
- It is very widely used
 - It was originally developed by Sun Microsystems (now bought by Oracle), but is now implemented in most (all?) Unix systems
 - Also implemented on Windows
 - The NFS distributed file system is based on Sun RPC
- It is **platform independent**
 - A Linux computer can call a procedure on a Solaris box, etc.
- But **not language independent**
 - Designed to call remote **C procedures**
 - But you can also use other languages that have gateways to C

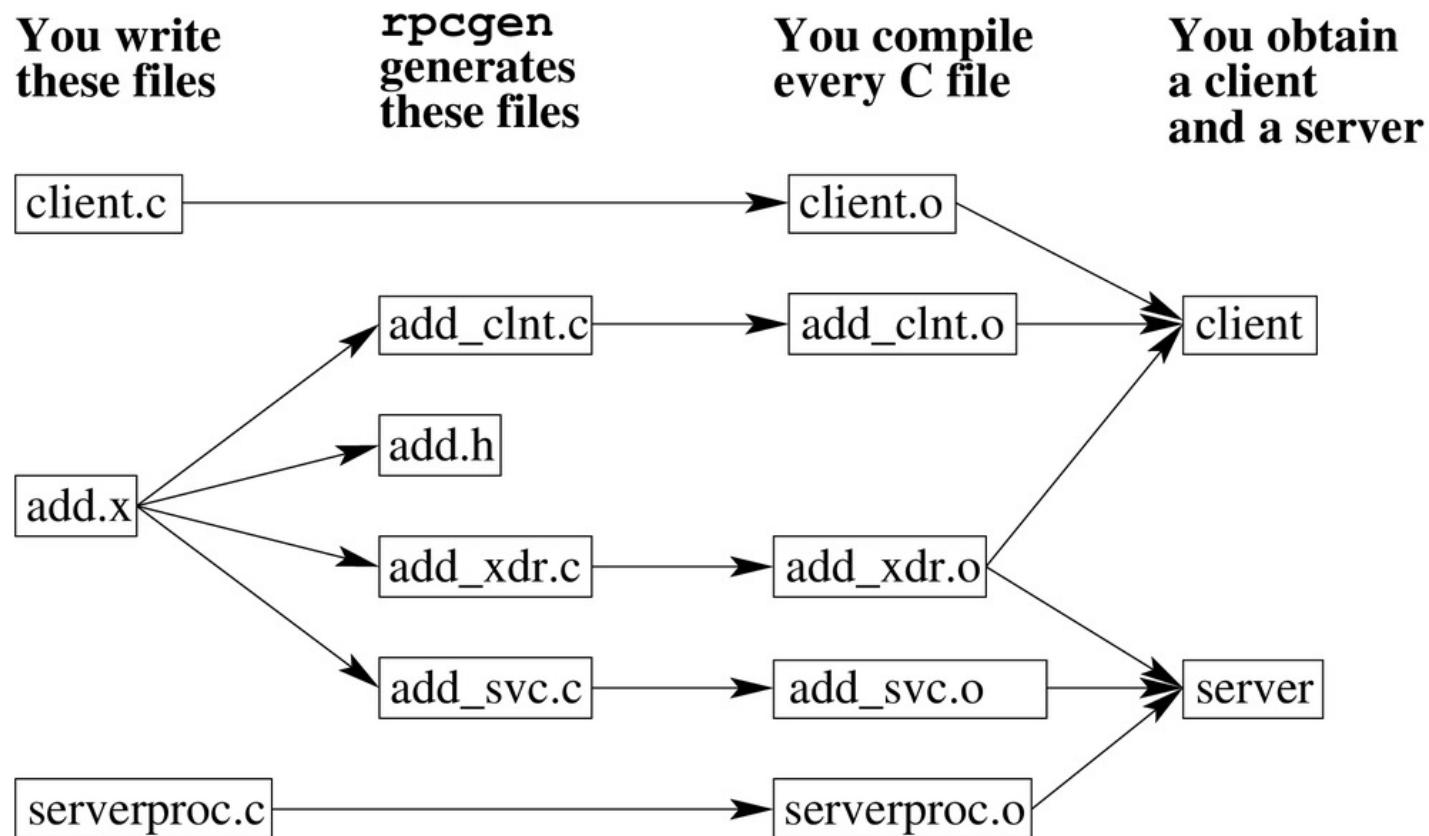
Writing an RPC Program [1/3]

- To write a minimalist RPC program, you must write:
 - A C procedure to be remotely called: **remoteproc.c**
 - A specification of the procedure: **remoteproc.x**
 - A client program that calls the procedure: **client.c**

Writing an RPC Program [2/3]

- Based on **remoteproc.x**, the program rpcgen generates:
 - A **header file** that you will include in both programs: **remoteproc.h**
 - A **client stub**, that your client can use to send an RPC:
remoteproc_clnt.c
 - A **server skeleton**, which will call your procedure when a request is received: **remoteproc_svc.c**
 - **Internal functions** to convert the procedure parameters into network messages and vice-versa: **remoteproc_xdr.c**
- Beware: on Solaris, rpcgen generates K&R code by default.
 - Use **rpcgen -C** to generate ANSI code

Writing an RPC Program [3/3]



A Bit of Terminology

- One computer can be a server of multiple procedures
- But how is it organized?
 - A server may host several **programs** (identified by a **program number**)
 - Each program may have several subsequent **versions** (identified by a **version number**)
 - Each version of a program may contain one or more **procedures** (identified by a **procedure number**)
- Program numbers are 32-bit hexadecimal values (e.g., 0x20000001)
 - As a user, you can choose any program number between 0x20000000 and 0x3FFFFFFF
 - **But make sure program numbers are unique!**
 - You cannot have several programs with the same number on the same machine
- Version and procedure numbers are integers (1, 2, ...)

An RPC Example

- Start by writing the specification file: **add.x**

```
struct add_in {      /* The arguments of the procedure */
    long arg1;
    long arg2;
};

typedef long add_out; /* The return value of the procedure */

program ADD_PROG {
    version ADD_VERS {
        add_out ADD_PROC(add_in) = 1; /* Procedure number = 1 */
        } = 1;                      /* Version number = 1 */
    } = 0x20001234;                /* Program number = 0x20001234 */
```

An RPC Example

- This file contains specifications of:
 - A structure **add_in** containing the arguments
 - A typedef **add_out** containing the return values
 - A program named **ADD_PROG** whose number is 0x20001234
 - The program contains one version with value **ADD_VERS** = 1
 - The version contains one procedure with value **ADD_PROC** = 1
 - This procedure takes an **add_in** as parameter, and returns an **add_out**

- Remarks
 - Your procedures can only take one input argument and return one output return value
 - If you need more arguments or return values, group them into a structure (like **add_in**)
 - Fields that are represented as pointers (e.g., `char *`) should not be assigned `NULL`! They should point to some actual content.

An RPC Example

- Generate stubs:

```
rpcgen add.x
```

- **add.h** contains various declarations:

```
#define ADD_PROG 0x20001234          /* Program nb */
#define ADD_VERS 1                      /* Version nb */
#define ADD_PROC 1                      /* Procedure nb */
add_out * add_proc_1(add_in *, CLIENT *);
add_out * add_proc_1_svc(add_in *, struct svc_req *);
```

- **add_proc_1**: the stub, i.e., the procedure that the client program will call
- **add_proc_1_svc**: the actual procedure that you will write and run at the server

- **add_clnt.c** contains the implementation of `add_proc_1`
- **add_svc.c** contains a program which calls your procedure `add_proc_1_svc` when it receives a request
- **add_xdr.c**: marshall/unmarshall routines

An RPC Example

- Write your server procedure: **serverproc.c**

```
#include "add.h"

add_out *add_proc_1_svc(add_in *in, struct svc_req *rqstp) {
    static add_out out;
    out = in->arg1 + in->arg2;
    return(&out);
}
```

- **rqstp** (request pointer): contains some information about the requester
 - Its IP address, etc.

An RPC Example

- Compile your server
- You need to compile together your procedure, the (generated) server program, the (generated) marshal/unmarshal procedures and the *nsl* library
 - The ns_l library contains the RPC runtime

```
$ gcc -c serverproc.c
$ gcc -c add_svc.c
$ gcc -c add_xdr.c
$ gcc -o server serverproc.o add_svc.o add_xdr.o -lnsl
```

- You can start your server:

```
./server
```

An RPC Example

- Write the client program: client.c

```
#include "add.h"
int main(int argc, char **argv) {
    CLIENT *cl;
    add_in in;
    add_out *out;

    if (argc!=4) {
        printf("Usage: client <machine> <int1> <int2>\n\n";
        return 1;
    }

    cl = clnt_create(argv[1], ADD_PROG, ADD_VERS, "tcp");
    in.arg1 = atol(argv[2]);
    in.arg2 = atol(argv[3]);
    out = add_proc_1(&in, cl);
    if (out==NULL) { printf("Error: %s\n",clnt_sperror(cl,argv[1])); }
    else { printf("We received the result: %ld\n",*out); }
    clnt_destroy(cl);
    return 0;
}
```

An RPC Example

- You must first create a client structure using **clnt_create()**

```
#include <rpc/rpc.h>
CLIENT *clnt_create(char *host, u_long prog, u_long vers, char *proto);
```

- **host**: the name of the server machine
 - **prog, vers**: the program and version numbers
 - **proto**: the transport protocol to use (“tcp”, or “udp”)
-
- Then you can call the (generated) client procedure **add_proc_1()** to send the RPC
 - When you are finished, you destroy the client structure
 - A client structure can be used multiple times without being destroyed and re-created

An RPC Example

□ Compile your client

```
$ gcc -c client.c
$ gcc -c add_clnt.c
$ gcc -c add_xdr.c
$ gcc -o client client.o add_clnt.o add_xdr.o -lnsl
```

□ Try it all

- Start your server

```
$ ./server
```

- Send a request

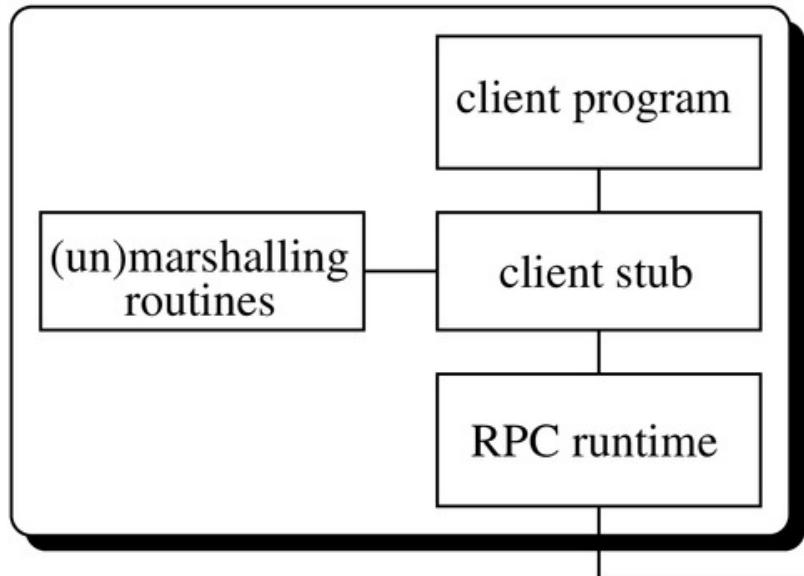
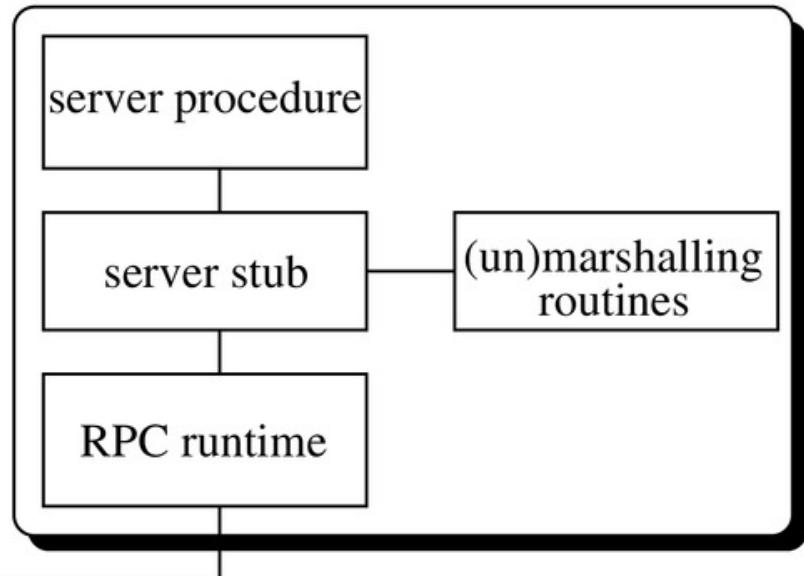
```
$ ./client flits.cs.vu.nl 30 7
We received the result: 37
```

Concurrent RPC Servers

- By default, generated RPC servers are iterative
 - **QUESTION:** how can you check that by yourselves?
- Some versions of **rpcgen** allow one to generate server code which implements one-thread-per-client
 - Solaris has a multi-thread extension, Linux doesn't
 - On Solaris: use `rpcgen -A -M` instead of `rpcgen` to generate a multithreaded server
 - You also need to change your client program and server procedure
- **QUESTION:** Can you transform an iterative RPC server into a concurrent one by adding (for example) `fork()` calls in your server procedure?

Inside Sun RPC

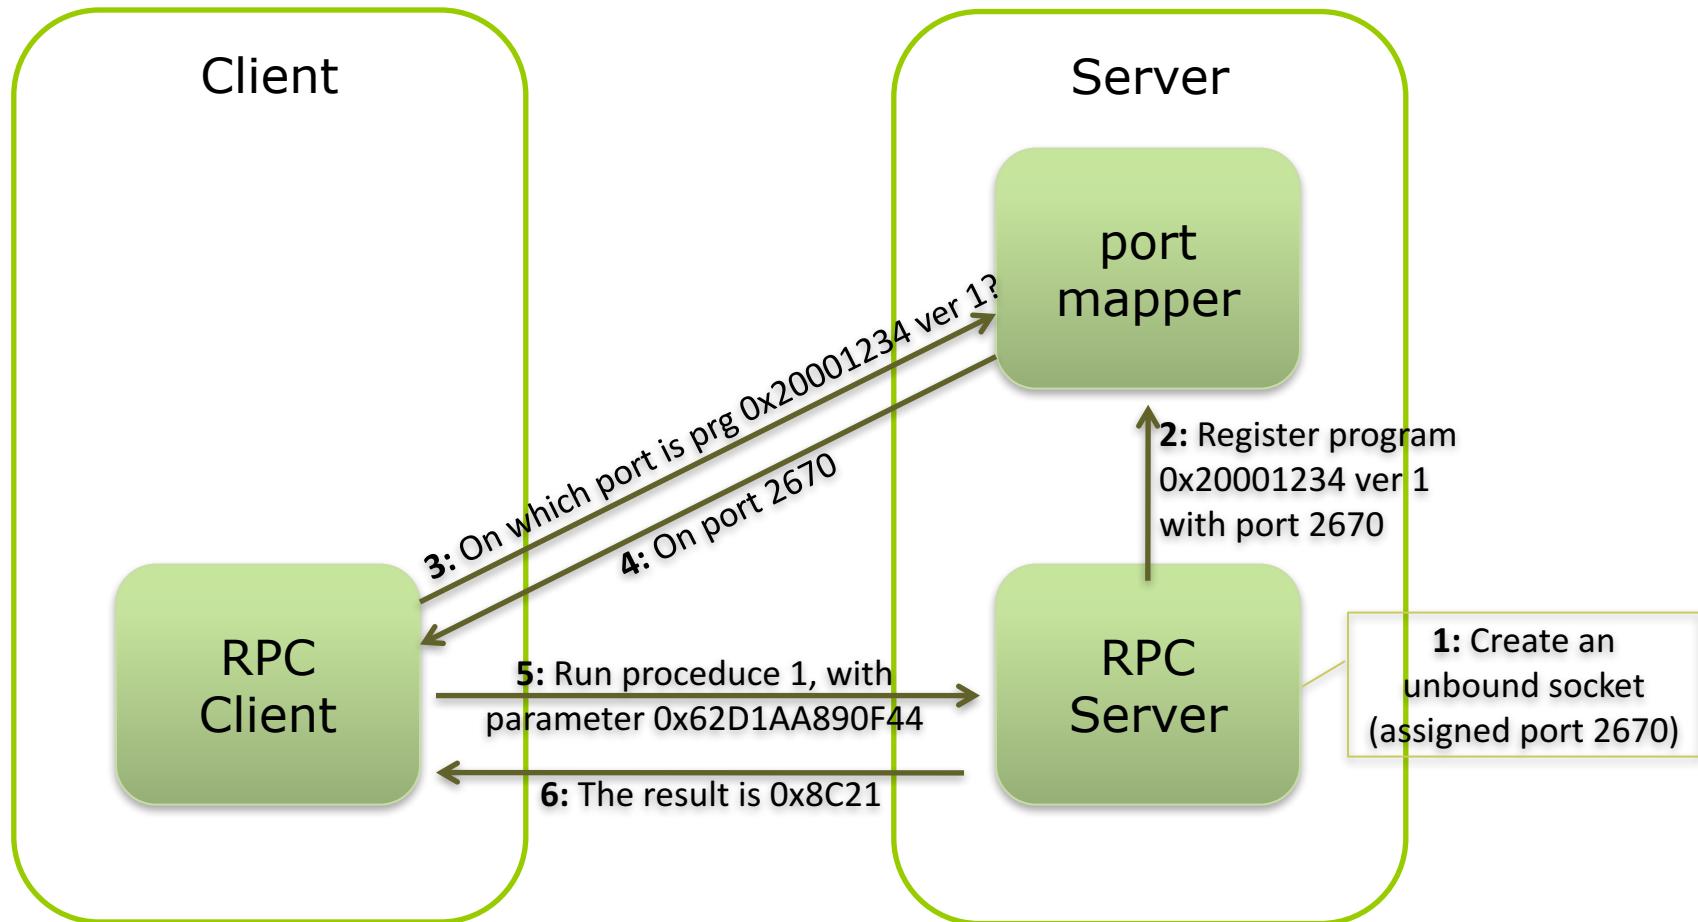
The Global Picture

Client**Server**

The Port Mapper [1/3]

- Did you notice that we did not specify a port number for our *add* server?
 - This could be done by using a well-known port number for all RPCs
 - But in that case, we would not be able to run multiple servers in the same machine simultaneously
- The RPC solution: **The Port Mapper**
 - The port mapper is a server running on **port 111** (well-known port)
 - When your RPC server starts, it does not bind its socket to a specific port. Instead, it registers whatever port has been assigned to it by the OS to the local port mapper
 - When you create a client with **clnt_create**, you automatically contact the remote port mapper to ask for the port number of the server you want to contact

The Port Mapper [2/3]



The Port Mapper [3/3]

- You can also query a port mapper by hand:

```
$ rpcinfo -p acropolis.cs.vu.nl
program vers proto  port
 100000    2   tcp    111  portmapper
 100000    2   udp    111  portmapper
 55848960    1   udp   32898
 55848960    1   tcp   58411
$
```

A Glimpse of the Protocol

- The RPC runtime does not know anything about **procedure names** or **prototypes**
- All it does is manage **requests** containing:
 - A host name
 - A program number
 - A version number
 - A procedure number
 - A buffer containing the arguments
- And **responses** containing:
 - A buffer with the return values
- The RPC runtime does not know (nor does it care) what the procedures are doing or what kind of arguments they take
 - It just ships buffers

Transport Issues

- RPC invocation can be carried out either by **UDP** or **TCP**
- **UDP** is **lighter, faster**
 - But what happens if the **request gets lost**?
 - After a timeout, the client system automatically resends the request, until it receives a response
 - Now, what happens if the **response gets lost**?
 - The server has been invoked twice
 - **UDP-based RPC** implements the **At-Least-Once semantics**: each invocation is performed on the server at least once, but sometimes several times
- You must be aware of these issues
 - If you choose **UDP**, then your server must be **idempotent** (i.e., applying the same operation twice has the same effect as applying it once)
 - If your server is not idempotent, then you **must** use **TCP**
 - **QUESTION:** Why would UDP not work?
 - **QUESTION:** How could we make it work for non-idempotent services?

(Un)Marshalling Parameters

- Parameters must be converted from variables into buffers and vice-versa
 - The lowest level only knows about parameters in the form of buffers to be shipped
 - The user programs only know about parameters in the form of C variables
 - The XDR (**eXternal Data Representation**) layer is in charge of the conversion
- The XDR buffer representations must be architecture independent
 - To avoid problems when a Little-Endian client sends requests to a Big-Endian server, or vice versa

(Un)Marshalling Parameters

- The XDR library contains basic encoding/decoding functions
- Example: to (un)marshal an integer

```
#include <rpc/xdr.h>
bool_t xdr_int(XDR *xdrs, int *ip);
```

- The xdrs structure contains a file descriptor of the socket where to read/write the buffers
- It also contains a flag to indicate whether xdr_int must encode or decode the buffer
 - Encode: convert *ip into a message buffer and write it to the socket
 - Decode: read a message buffer from the socket and convert it to *ip

(Un)Marshalling Parameters

- rpcgen generates functions to (un)marshal the input and output parameters of your RPC procedures:

```
#include "add.h"

bool_t xdr_add_in(XDR *xdrs, add_in *objp) /* to convert add_in */
{
    if (!xdr_long(xdrs, &objp->arg1)) return FALSE;
    if (!xdr_long(xdrs, &objp->arg2)) return FALSE;
    return TRUE;
}

bool_t xdr_add_out(XDR *xdrs, add_out *objp) /* to convert add_out */
{
    if (!xdr_long(xdrs, objp)) return FALSE;
    return TRUE;
}
```

Client Stub

- The client stub generated by rpcgen is in charge of providing a simple interface to the client programs:
 - It calls the XDR functions to marshal the procedure parameters and to unmarshal the return values
 - It ships them to the low-level RPC layers to be transferred to the server

```
add_out *add_proc_1(add_in *argp, CLIENT *clnt) {
    static add_out clnt_res;

    memset((char *)&clnt_res, 0, sizeof(clnt_res));
    if (clnt_call(clnt, ADD_PROC,
                  (xdrproc_t) xdr_add_in, (caddr_t) argp,
                  (xdrproc_t) xdr_add_out, (caddr_t) &clnt_res,
                  TIMEOUT) != RPC_SUCCESS) {
        return (NULL);
    }
    return (&clnt_res);
}
```

Server Skeleton

- The server skeleton generated by **rpcgen** does the following:
 1. Defines a function to reply to any request directed at it (same program and version number, any procedure number)
 - It checks the procedure number of the incoming request
 - Calls the appropriate function to unmarshal the incoming parameters
 - Calls the user procedure
 - Marshals the return values
 2. Creates a **UDP and TCP** server socket, and registers them to the port mapper
 3. Registers the function defined in step 1 to the RPC system:
 - Program and version numbers
 - The UDP and TCP ports
 - The procedure to serve incoming requests
 4. Waits for RPC requests to arrive

CORBA

(not included in the exam material, read optionally)

CORBA

- CORBA
 - Common Object Request Broker Architecture
 - A remote object invocation system
- Language-independent, Architecture independent
 - e.g., a C++ object on a Linux computer can invoke a Java object on a Windows computer
- Corba is a specification (not a *piece of software*)
 - Corba defines the protocols, the interfaces, etc.
 - Any vendor can implement the specification
 - An implementation is called an ORB (**O**bject **R**equest **B**roker)
 - Each ORB can contain additional tools to help developers
 - Some are standardized, some not
 - There are many ORBs
 - For instance, the JDK-1.2 contains an ORB

The Interface Definition Language

- Since CORBA is language-independent, you must represent server interfaces in a way that every language can understand: the **Interface Definition Language (IDL)**
- Example:

```
module Stock {
    exception Invalid_Stock {};

    struct Info {
        long high;
        long low;
        long last;
    };

    interface Quoter {
        Info get_quote(in string stock_names) raises(Invalid_Stock);
    };
}
```

IDL Compilers

- Each ORB provides an **IDL compiler**
 - It compiles the interface specification into stubs and skeletons
 - Example: the JDK-1.2 ORB contains an IDL compiler called **idltojava**, or **idlj**

- The CORBA specification defines **language mappings**
 - How to represent an IDL interface in a given language
 - Two ORBs must generate exactly the same language specific interface from the same IDL
 - There are mappings for C, C++, Java, Python, Perl, Ruby, Ada, COBOL, Smalltalk, Objective-C, Lisp, etc.

Example: The Java Mapping

- The IDL concepts are translated into Java concepts:

IDL	Java
Module	Package
Interface	Interface
Operation	Method
struct	Object
etc.	

Implementing a Server and a Client

- Once you have compiled your IDL interface, you can fill in the implementation for your server object
 - Details vary by language (of course)
- To write a server program
 1. Initialize the ORB
 2. Instantiate (at least) one server object
 3. Connect the server object to the ORB
 4. Wait for requests
- To write a client program
 1. Initialize the ORB
 2. Obtain a reference to a server object
 3. Invoke the method

Cross-language Remote Invocations

- To send requests from a client written in language A to a server written in language B:
 1. Write an IDL specification of the server
 2. Compile the specification to language B, implement the server in language B
 3. Compile the specification to language A, implement the client in language A
 4. Run it!

Object Adapters

- An **object adapter** is in charge of managing server objects
 - Can server objects be invoked concurrently? Each object can decide which policy it prefers
 - Sometimes you want to store server objects on disk, and instantiate them only when a request is directed to it
 - etc.

- The **Portable Object Adapter (POA)** is a standardized interface to such an Object Adapter

The Dynamic Invocation Interface

- Sometimes you don't know the exact interface of the server at compile time
- The **Dynamic Invocation Interface (DII)** allows you to dynamically construct a request before invoking it
- Beware: this code probably doesn't work

```
org.omg.CORBA.Object obj = ...;
org.omg.CORBA.Request req = obj->getRequest("set_quote");
req.addInArg("VU");
req.addInArg(42);
req.setReturnType(new org.omg.CORBA.int());
req->invoke();
int result = req.return_value();
```

Corba Services

- In addition to basic remote invocation, CORBA defines services:
 - A CORBA service is a complete server object which is provided by an ORB
 - They provide high-level functions to help you develop applications

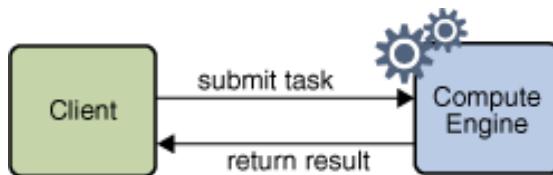
- Examples
 - The Naming Service associates a name with an object
 - The Trading Service allows you to search for an object **by functionality**
 - The Event Service allows you to be notified when certain events happen
 - The Interface Repository allows you to ask “What is the interface of this object?”

RMI

Remote Method Invocation

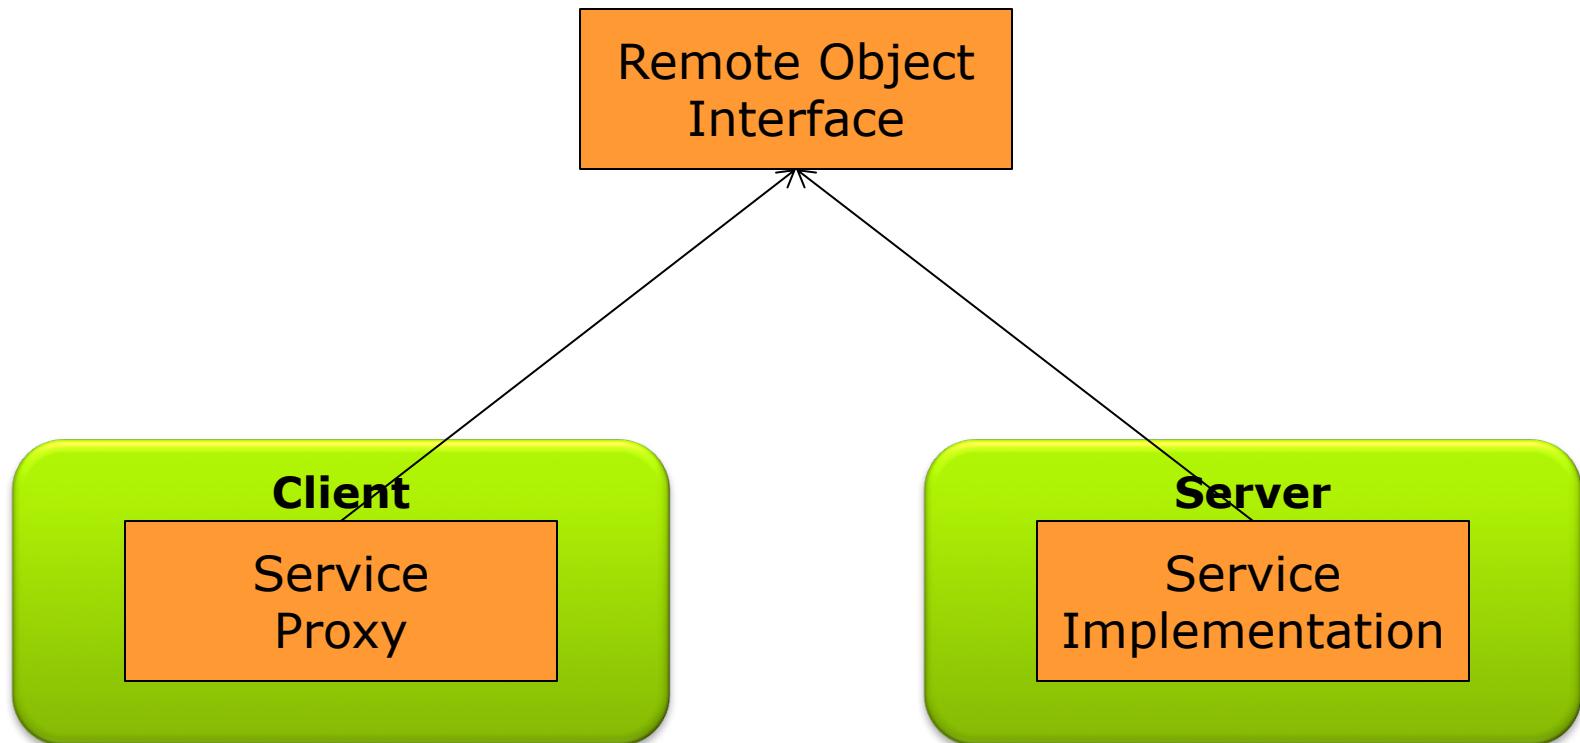
Java RMI

- RMI (Remote Method Invocation)
 - allows Java programs to invoke methods of remote objects

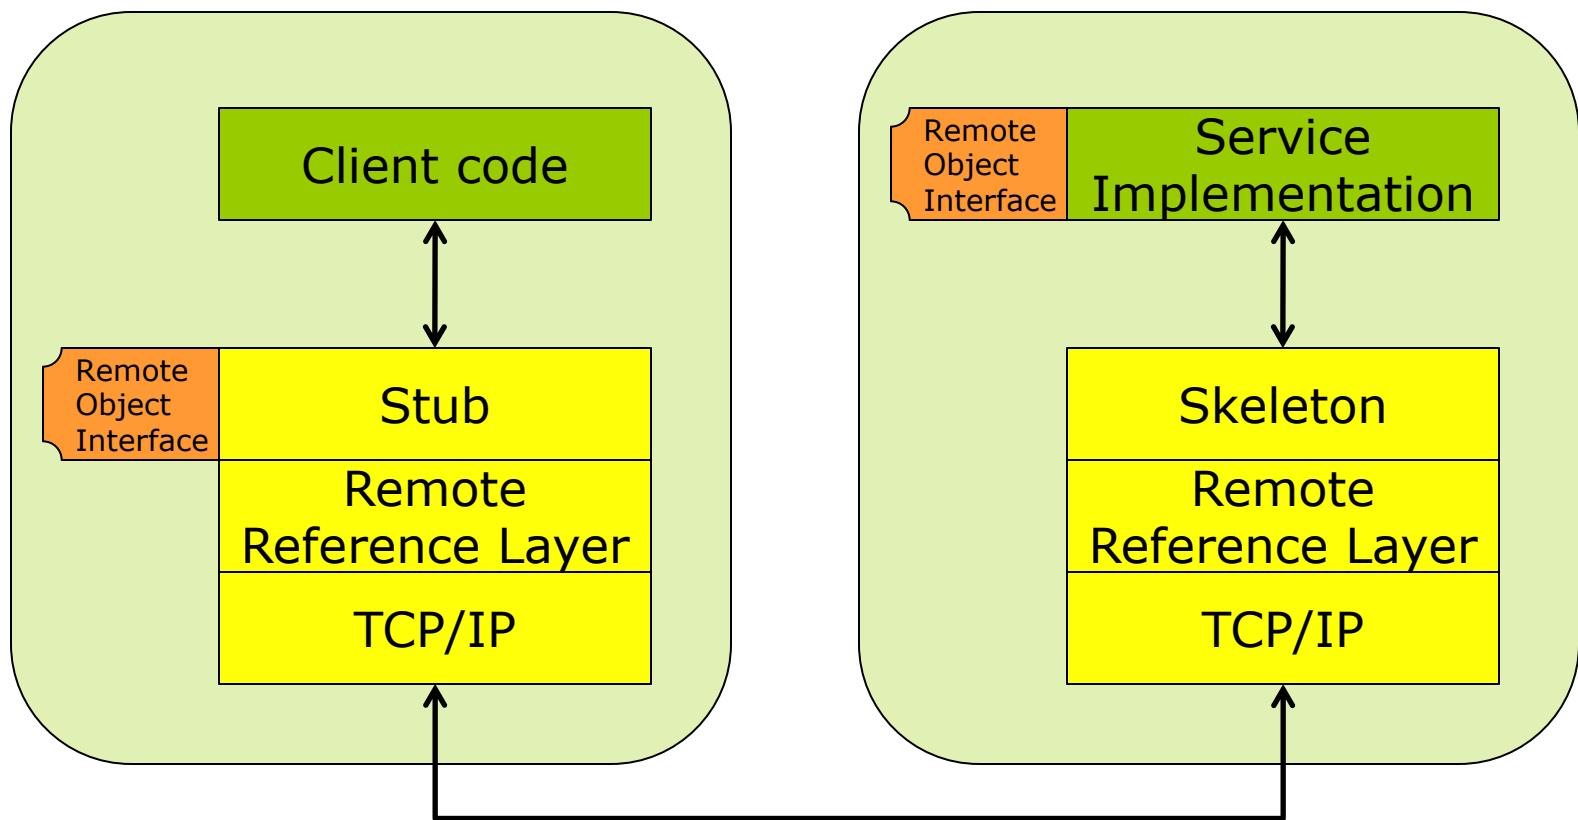


- Access distributed objects (almost) identically to local objects
 - Same syntax (arguments, return values, etc.)
 - Same semantics (exceptions)
- Scope
 - Only in Java
 - Significant changes in Java 1.2 (incompatible to earlier RMI)

RMI Interface



RMI Architecture



Transparency

- To the client, a remote object appears exactly like a local object
 - Except that **you must bind it first!**
- This is possible thanks to **interfaces**
 - You write the **interface** to the remote object
 - You write the **implementation** for the remote object
 - RMI automatically creates a **stub class** which implements the remote object interface
 - The client accesses the stub exactly the same way it would access a local copy of the remote object

The RMI Registry

- Java RMI needs a naming service (like Sun RPC's port mapper)
 - Servers register contact address information
 - Clients can locate servers
- This is called **RMI Registry**
- Unlike in RPC, you must start the RMI Registry yourself
 - **Question:** Where?
 - on each machine that hosts **server objects**
 - It is a program called **rmiregistry**
 - By default runs on port 1099 (but you can specify another port number: **rmiregistry <port_nb>**)
- Programs access the registry via the **java.rmi.Naming** class

An RMI Example

- To write a minimalist RMI program you must write:
 - An interface for the remote object: **Remote.java**
 - An implementation for the remote object: **RemoteImpl.java**
 - A server which will run the remote object: **RemoteServer.java**
 - A client to access the server: **RemoteClient.java**

- The RMI compiler **rmic** will generate the rest:
 - For Java version up to 1.1
 - A client stub: **RemoteImpl_Stub.class** (already compiled)
 - A server skeleton: **RemoteImpl_Skel.class** (already compiled)
 - For Java version 1.2 ... 1.4 (Java 2 ... Java 4)
 - A single stub, used for both client and server: **RemoteImpl_Stub.class**
 - For Java 1.5 (Java 5) and higher
 - Nothing is needed!
 - Stubs and skeletons are included in the Java library

An RMI Example (continued)

- Start by writing the interface for the remote object: **Calculator.java**

```
public interface Calculator extends java.rmi.Remote {  
    public long add(long a, long b)  
        throws java.rmi.RemoteException;  
  
    public long sub(long a, long b)  
        throws java.rmi.RemoteException;  
}
```

- You must respect a few rules
 - The interface must extend the **java.rmi.Remote** interface
 - All methods must throw the **java.rmi.RemoteException** exception
- Compile the interface

```
$ javac Calculator.java
```

An RMI Example (continued)

- Write an implementation for the remote object: CalculatorImpl.java

```
public class CalculatorImpl
    extends java.rmi.server.UnicastRemoteObject
    implements Calculator {

    // Implementations must have an explicit constructor
    public CalculatorImpl() throws java.rmi.RemoteException {
        super();
    }

    public long add(long a, long b) throws java.rmi.RemoteException {
        return a + b;
    }

    public long sub(long a, long b) throws java.rmi.RemoteException {
        return a - b;
    }
}
```

An RMI Example (continued)

- The implementation class must respect a few constraints
 - It must implement the interface (of course!)
 - It must inherit from the **java.rmi.server.UnicastRemoteObject** class
 - It must have an explicit constructor which throws the **java.rmi.RemoteException** exception
- Compile the implementation class

```
$ javac CalculatorImpl.java
```

An RMI Example (continued)

- If you are using Java version earlier than 5.0, generate the stub and skeleton:

```
$ rmic CalculatorImpl
```

- This generates directly the **CalculatorImpl_Stub.class** and (for Java < 2.0) the **CalculatorImpl_Skel.class** files
 - Already compiled to byte code, you are not expected to run javac

An RMI Example (continued)

- Write a server program: CalculatorServer.java

```
import java.rmi.Naming;

public class CalculatorServer {
    public CalculatorServer() {
        try {
            Calculator c = new CalculatorImpl();
            Naming.rebind("rmi://localhost/CalculatorService", c);
        } catch (Exception e) {
            System.out.println("Trouble: " + e);
        }
    }

    public static void main(String args[]) {
        new CalculatorServer();
    }
}
```

An RMI Example (continued)

- The server program creates a **CalculatorImpl** object
- It registers the object to the local RMI registry

```
rebind(String name, Remote obj)
```

- Associates a **name** to an **object**
- Names are in URL form: **rmi://<host_name>[:port]/<service_name>**

- The server will wait for incoming requests
- Compile your server

```
$ javac CalculatorServer.java
```

An RMI Example (continued)

- Write a client program: CalculatorClient.java

```
import java.net.MalformedURLException;
import java.rmi.Naming;

public class CalculatorClient {
    public static void main(String[] args) {
        try {
            Calculator c = (Calculator)Naming.lookup(
                "rmi://flits.few.vu.nl/CalculatorService");
            System.out.println( c.add(4, 5) );
            System.out.println( c.sub(4, 3) );
        }
        catch (Exception e) {
            System.out.println("Received Exception:");
            System.out.println(e);
        }
    }
}
```

An RMI Example (continued)

- Before invoking the server, the client must lookup the registry
 - It must provide the URL for the remote service
 - It gets back a stub which has exactly the same interface as the server
 - It can use it as a local object: **long x = c.add(4,5)**

- Compile your client:

```
$ javac CalculatorClient.java
```

An RMI Example (end!)

□ Try your program!!

- Start the RMI Registry: rmiregistry

```
$ rmiregistry &
```

- The registry must have access to your classes
- Either start the registry in the same directory as the classes, or make sure the directory is listed in the \$CLASSPATH variable

- Start your server

```
$ java CalculatorServer
```

- Start your client

```
$ java CalculatorClient  
9  
1  
$
```

Using RMI in a Distributed Context

- First, try your RMI program on a single host
 - It will work only if there is an operational TCP configuration
- To use the client and server on different hosts, you must provide the right files to the right entities:
 - The **server** and the **rmiregistry** need the following files:
 - **Calculator.class**: the remote object interface
 - **CalculatorImpl.class**: the server object implementation
 - **CalculatorServer.class**: the server program
 - The **client** needs the following files:
 - **Calculator.class**: the remote object interface
 - **CalculatorClient.class**: the client program

Passing Parameters

- Depending on the parameter's type, a different method is followed
- **Primitive types** (e.g., int, float, long, char, boolean) are copied by value
- **Object types** are serialized and transferred
 - The object itself, plus every other object that it refers to (recursively)
 - You must be careful:
 - Remote invocations **pass objects by value** (whereas local invocations pass objects **by reference**)
 - It is very easy to transfer huge quantities of data without noticing
 - (e.g., you have a whole database in memory, and you transfer an object that contains a reference to the database)
- **Remote Objects** (i.e., inheriting from java.rmi.Object) are not transferred!
 - Instead, a distributed reference to this object is transferred
 - Any invocation to this object will result in an RMI request

Questions

- Can you also operate RMI over UDP?
 - No! Only TCP.

- Why does the main thread in CalculatorServer.java end?
 - A Java program terminates when all threads have terminated, unlike a C program that terminates when the main thread terminates

- bind() vs. rebind() ?

- When do we get exceptions?
 - Exceptions raised on the server are propagated and thrown on the client

Questions

- Is RMI blocking on the client side?
 - YES: The client blocks until the server replies

- How can we make RMI non-blocking for the client's side?
 - **Option 1:** Make separate thread and let it block
 - **Option 2:** Or make the client also a Remote Object, pass it on to the server, and let it call you back once it's done

- Is the server iterative?
 - No!
 - Two or more requests can be handled simultaneously

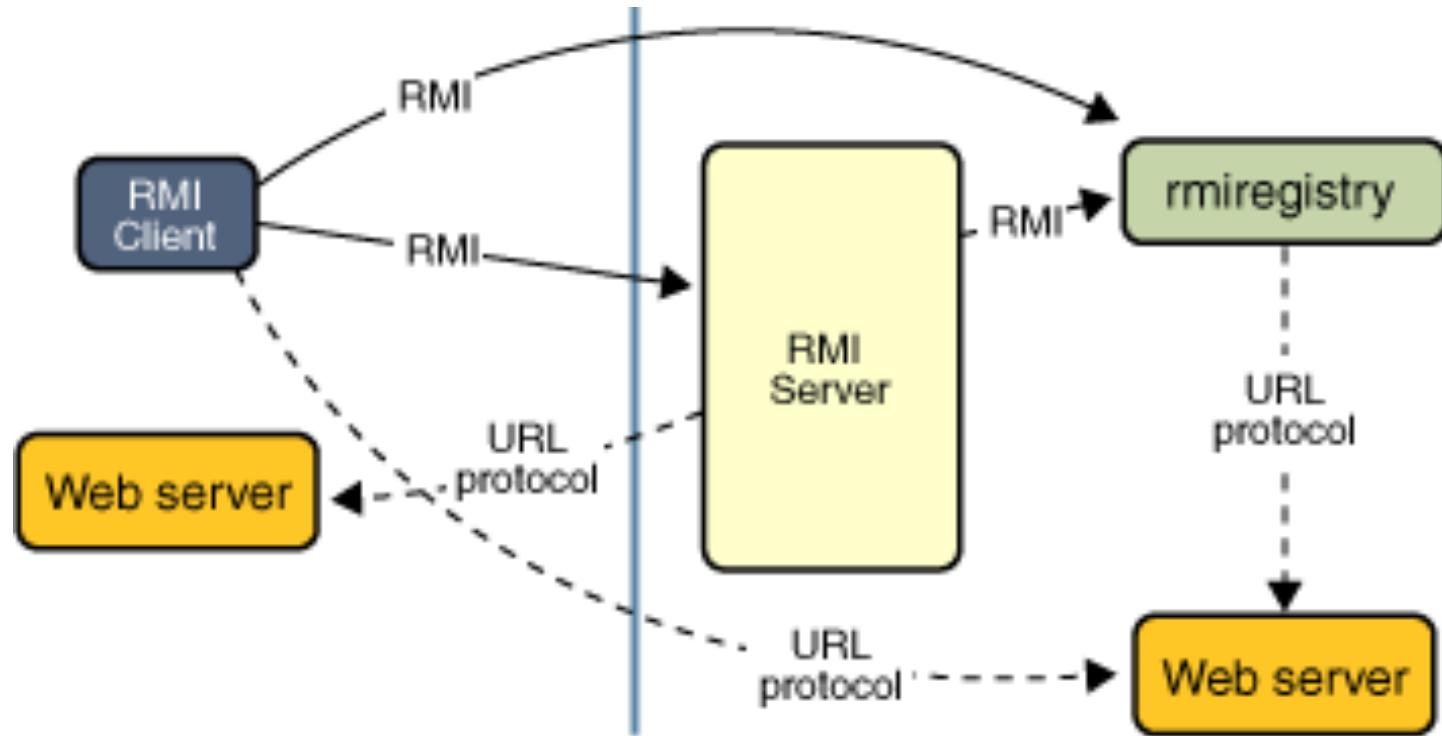
Remote Codebase [1/2]

- Why does the Registry require access to the implementation files?
 - Because a JVM can send to another JVM a (serializable) object whose class is not known by the receiver (although its interface *is* known)

- What if we are using a remote registry? (i.e., running on a different machine than the one hosting the remote object)
 - You can specify a remote codebase for Foo like this:

```
$ java -Djava.rmi.server.codebase=http://myhost/~abc/myclasses/ Foo
```

Remote Codebase [2/2]



Security

- When a Java application needs to download code of a serializable class from a remote location, security has to be considered
- To set your custom security policy, start your application like this:

```
$ java -Djava.security.policy=myPolicy.txt Foo
```

- If your program crashes with a security exception, try it out (temporarily) with a policy file like this:

```
grant {  
    // Allow everything for now  
    permission java.security.AllPermission;  
};
```

- The default policy file is found at
<JAVAHOME>/lib/security/java.policy

Garbage Collection

- Garbage Collection is hard in a single machine...
 - ...very hard in a distributed setting!
- The server keeps a reference counter on
 - local references for a Remote Object
 - client references for a Remote Object
- The server garbage collects Remote Objects that are not referenced locally nor remotely
- A Remote Object can implement the **java.rmi.server.Unreferenced** interface to get a notification via the **unreferenced()** method when all clients have dropped it

Application Architecture Example

- Discussion in class:
 - How would you implement a Multiplayer Chess system by RMI?

Remote Codebase [2/2]

